

# CS244b – Distributed Systems

**Instructor:** David Mazières

**CAs:** Hiroshi Mendoza and Seo Jin Park

Stanford University

# Outline

- 1 Administrivia
- 2 Remote procedure call
- 3 Consensus in asynchronous systems

# Administrivia

- **Class web page:** <http://cs244b.scs.stanford.edu/>
- **All handouts and lecture notes on line**
  - Please print them out yourselves
- **Each class will involve discussing papers**
  - Print, read the papers before class
  - Class (except SCPD) and mailing list participation counts for grade
  - We will post discussion notes afterwards
- **Homework questions before most classes (see [syllabus](#))**
  - Turn in on paper at start of class (SCPD can submit remotely)
- **Staff mailing list:** [cs244b-staff@scs.stanford.edu](mailto:cs244b-staff@scs.stanford.edu)
  - Please email all staff rather than individual members

# Programming assignments

- **Two solo programming assignments in C++11 (or later)**
  - Goal: familiarize you with RPC, consensus, consistency
- **Final project**
  - Perform a small research project in teams of 1–3 students
  - Welcome to use code from first labs
  - Use ideas from papers we've discussed in class
  - Turn in short paper, make presentation
- **Presentations: Monday, December 11, (12:30–6:30pm??)**
  - Present project in mini-conference
  - We will serve food
  - Might need second slot, student PC, or parallel tracks given enrollment

# Grading

- **Grading based on four factors:**
  1. Class participation and homework questions ( $c$ , where  $0 \leq c \leq 1$ )
  2. Midterm and final quizzes ( $q$ )
  3. Lab assignments ( $l$ )
  4. Final project paper & presentation ( $p$ )
- **Combined as follows (subject to adjustment):**
  - Compute average:  $a = q/4 + l/4 + p/2$ .
  - Adjusted score is: **if**  $p > a$  **then**  $cp + (1 - c)a$  **else**  $a$ .
- **Final project is most important component**
- **With participation, good project overrides bad quiz/lab**

# Why study distributed systems?



- **Most real systems are actually distributed systems**
- **If you want fault-tolerance or scalability**
  - Must replicate across multiple machines
- **If you want systems that span administrative realms**
  - Web sites, peer-to-peer systems, communication systems

# Class topics

- Distributed programming models
- Dealing with failure, including Byzantine failure
- Scalability
- Techniques: Consensus, Replication, Consistency...
- Case studies: production systems at Google, Amazon, ...
- Byzantine-agreement-based Blockchain mechanisms

# Outline

- 1 Administrivia
- 2 Remote procedure call
- 3 Consensus in asynchronous systems



# Remote procedure call (RPC)

- **Procedure calls are a well-understood mechanism**
  - Transfer control and data on single computer
- **RPC's goal is to make distributed programming look like as much as possible like normal programming**
  - Code libraries provide APIs to access functionality
  - RPC servers export interfaces accessible through local APIs
  - See [\[Birrell\]](#) for good description of one implementation
- **Implement RPC through request-response protocol**
  - Procedure call generates network request to server
  - Server return generates response
- **Good example of how distributed systems differ...**

# Procedure vs. RPC

- **Consider the following ordinary procedure:**

```
bool add_user(string user, string password);
```

- **Possible return values:** `true`, `false`
- **Now say you have an RPC version**
  - Must somehow set up connections, bind to server, think about authentication, etc., but ignore all that for now
- **What are the possible return values of `add_user` RPC?**

# Procedure vs. RPC

- Consider the following ordinary procedure:

```
bool add_user(string user, string password);
```

- Possible return values: `true`, `false`
- Now say you have an RPC version
  - Must somehow set up connections, bind to server, think about authentication, etc., but ignore all that for now
- What are the possible return values of `add_user` RPC?
  1. `true`
  2. `false`
  3. “I don’t know”

# RPC Failure

- **Normal procedure call has fate sharing**
  - Single process: if callee fails, caller fails, too
- **RPC introduces more failure modes**
  - Machine failures at only one end (caller/callee)
  - Communication failures
- **Result: RPCs can return “failure” instead of results**
- **What are the possible outcomes after failure?**
  - Procedure did not execute
  - Procedure executed once
  - Procedure executed multiple times
  - Procedure partially executed
- **Many systems aspire to “at most once semantics”**

# Implementing at most once semantics

- **Danger: Request message lost**
  - Client must retransmit requests when it gets no reply
- **Danger: Reply message may be lost**
  - Client may retransmit previously executed request
  - Okay if operations are idempotent, but many are not (e.g., process order, charge customer, ...)
  - Server must keep “replay cache” to reply to already executed requests
- **Danger: Server takes too long to execute procedure**
  - Client will retransmit request already in progress
  - Server must recognize duplicate—can reply “in progress”

# Server crashes

- **Danger: Server crashes and reply lost**
  - Can make replay cache persistent—slow
  - Can hope reboot takes long enough for all clients to fail
- **Danger: Server crashes during execution**
  - Can log enough to restart partial execution—slow and hard
  - Can hope reboot takes long enough for all clients to fail
- **Can use “cookies” to inform clients of crashes**
  - Server gives client cookie which is time of boot
  - Client includes cookie with RPC
  - After server crash, server will reject invalid cookie

# Parameter passing

- **Trivial for normal procedure calls**
- **RPC must worry about different data representations**
  - Big/little endian
  - Size of data types
- **RPC has no shared memory**
  - No global variables
  - How to pass pointers
  - How to garbage-collect distributed objects
- **How to pass unions over RPC?**

# Interface Definition Languages

- **Idea: Specify RPC call and return types in IDL**
- **Compile interface description with IDL compiler. Output:**
  - Native language types (e.g., C/Java/C++ structs/classes)
  - Code to *marshal* (serialize) native types into byte streams
  - *Stub* routines on client to forward requests to server
- **Stub routines handle communication details**
  - Helps maintain RPC transparency, but...
  - Still have to bind client to a particular server
  - Still need to worry about failures



# C++ RPC-related systems in use today

- XML or JSON over HTTP – no IDL, hard to parse
- **Cereal** – C++11 structure serializer
- Google **protobufs**, Apache **Thrift**
  - + Compact encoding, defensively coded (protobufs)
  - + Good support for incrementally evolving messages
  - Not complete system (protobufs), complex encoding, not C++11
- Apache **Avro** – self-describing messages contain schema
- **Cap'n Proto**, Google **FlatBuffers**
  - + Same representation in memory and on wire, very fast
  - Less mature, non-deterministic wire format, bigger attack surface
- **XDR (+ RPC)** – used by Internet standards such as **NFS**
  - + Simple, good features (unions, fixed- and variable-size arrays, ...)
  - Big endian, binary but rounds everything to multiple of 4 bytes

## Case study: XDR

```
enum MyEnum { NO, YES, MAYBE };

struct MyMessage {
    string name<16>;    /* up to 16 characters */
    string desc<>;      /* up to 232-1 characters */
    opaque cookie[8];   /* 8 bytes (fixed) */
    opaque sig<16>;     /* 0-16 bytes (variable-length) */
    unsigned int u;     /* Unsigned 32-bit integer */
    hyper ii;           /* Signed 64-bit integer */
    MyEnum me;          /* Another user-defined type */
    int ia[5];          /* Fixed-length array */
    int iv<>;           /* Variable length array */
    int iv1<5>;         /* Up to 5 ints */
    MyMessage *mep;     /* optional MyMessage (or NULL) */
};

typedef MyMessage *OptionalMyStruct;
```

# XDR base types

- All numeric values encoded in big-endian order
- `int`, `unsigned [int]`, all `enums`: 4 bytes
- `bool`: equivalent to “`enum bool { FALSE, TRUE }`”
- `hyper`, `unsigned hyper`: 8 bytes
- `float`, `double`, `quadruple`: 4-, 8-, or 16-byte floating point
- `opaque bytes[Len]` (fixed-size)
  - Encoded as content + 0–3 bytes padding to make size multiple of 4
- `string s<MaxLen>`, `opaque a<MaxLen>` (variable-size)
  - 4-byte length + content + (0–3 bytes) padding

# XDR containers and structs

- **(Fixed) arrays** – `MyType var[n]`
  - Encoded as  $n$  copies of `MyType`
- **Vectors** – `MyType var<>` **or** `MyType var<n>`
  - Can hold variable number  $(0-n)$  `MyTypes`
  - Encoded as 4-byte length followed by that many
  - Empty maximum length means maximum length  $2^{32} - 1$  `MyTypes`
- **Optional data** – `MyType *var`
  - Encoded exactly as `MyType var<1>`
  - Note this means single “present” bit consumes 4 bytes
- `struct` – **each field encoded in turn**

```
union type switch (simple_type which) {  
    case value_A:  
        type_A varA;  
    case value_B:  
        type_B varB;  
    /* ... */  
    default:  
        void;  
};
```

- **Must be discriminated, unlike C/C++**
- `simple_type` **must be** [unsigned] int, bool, **or** enum
- **Wire representation:**
  - 4-bytes for `which` + encoding of selected case
  - Special `void` type encoded as 0 bytes

# Outline

- 1 Administrivia
- 2 Remote procedure call
- 3 Consensus in asynchronous systems

# The identity function

- One of the simplest functions is the identity function
- E.g., in Haskell: `id x = x`
- In C++11:

```
template<typename T> inline T &&  
id(T &&t)  
{  
    return static_cast<T &&>(t);  
}
```

- The distributed equivalent turns out to be much harder
  - Problem: agents might not start with the same input
  - So to agree on output, must somehow pick one of the inputs

# Asynchronous systems<sup>1</sup>




- **A theoretical model for distributed systems**
  - Consists of a set of **agents** exchanging messages
  - No bound on message delays
  - No bound on the relative execution speed of agents
  - For convenience, model internal events such as timeouts as special messages, so the “network” controls all timing
- **Can't distinguish failed agent from slow network**
- **Idea of model is to be conservative**
  - Want robustness under any possible timing conditions
  - E.g., say backhoe tears fiber, takes a day to repair
  - Could see messages delays a *billion* times more than usual

---

<sup>1</sup>Unrelated to “asynchronous IO” as used in event-driven systems.

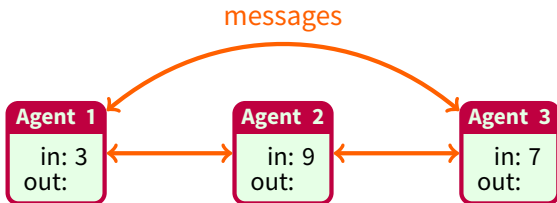


# The consensus problem



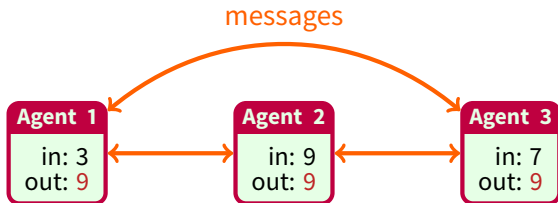
- **Goal:** For multiple agents to agree on an output value
- **Each agent starts with an input value**
  - Agents' inputs may differ; any agent's input is okay to output
- **Agents communicate following some *consensus protocol***
  - Use protocol to agree on one of the agent's input values
- **Once decided, agents output the chosen value**
  - Output is write-once (an agent cannot change its value)

# The consensus problem



- **Goal:** For multiple agents to agree on an output value
- **Each agent starts with an input value**
  - Agents' inputs may differ; any agent's input is okay to output
- **Agents communicate following some *consensus protocol***
  - Use protocol to agree on one of the agent's input values
- **Once decided, agents output the chosen value**
  - Output is write-once (an agent cannot change its value)

# The consensus problem



- **Goal:** For multiple agents to agree on an output value
  - **Each agent starts with an input value**
    - Agents' inputs may differ; any agent's input is okay to output
  - **Agents communicate following some *consensus protocol***
    - Use protocol to agree on one of the agent's input values
- **Once decided, agents output the chosen value**
- Output is write-once (an agent cannot change its value)

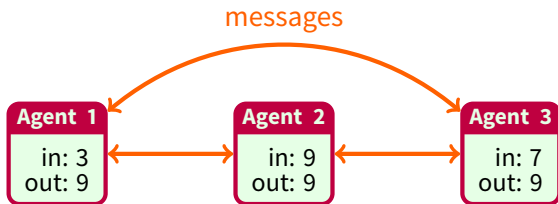
# Properties of a consensus protocol

- A consensus protocol provides *safety* if...
  - **Agreement** – All outputs produced have the same value, and
  - **Validity** – The output value equals one of the agents' inputs
- A consensus protocol provides *liveness* if...
  - **Termination** – Eventually non-failed agents output a value
- A consensus protocol provides *fault tolerance* if...
  - It can survive the failure of an agent at any point
  - *Fail-stop* protocols handle agent crashes
  - *Byzantine-fault-tolerant* protocols handle arbitrary agent behavior

## Theorem (FLP impossibility result)

*No deterministic consensus protocol provides all three of safety, liveness, and fault tolerance in an asynchronous system.*

# Bivalent states

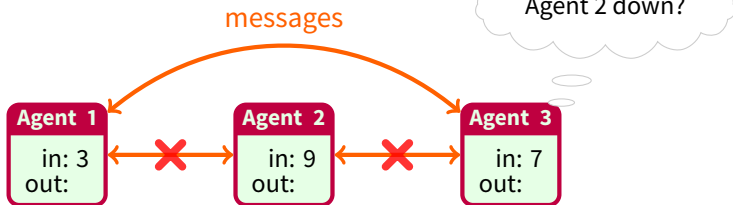


- Recall agents chose value 9 in last example
- But a network outage could look like agent 2 failing
  - If fault-tolerant, Agents 1 & 3 might decide to output 7
  - Once network back, Agent 2 must also output 7

## Definition (Bivalent)

An execution of a consensus protocol is in a **bivalent** state when the network can affect which value agents choose.

# Bivalent states

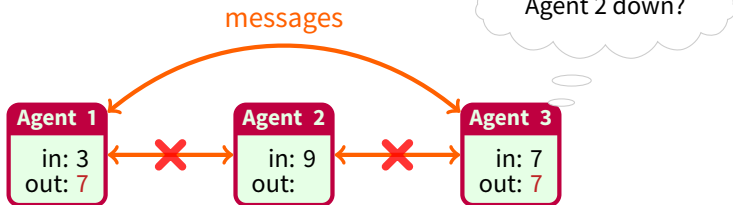


- Recall agents chose value 9 in last example
- But a network outage could look like agent 2 failing
- If fault-tolerant, Agents 1 & 3 might decide to output 7
  - Once network back, Agent 2 must also output 7

## Definition (Bivalent)

An execution of a consensus protocol is in a **bivalent** state when the network can affect which value agents choose.

# Bivalent states

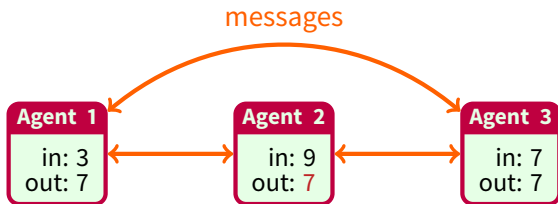


- Recall agents chose value 9 in last example
  - But a network outage could look like agent 2 failing
- If fault-tolerant, Agents 1 & 3 might decide to output 7
- Once network back, Agent 2 must also output 7

## Definition (Bivalent)

An execution of a consensus protocol is in a **bivalent** state when the network can affect which value agents choose.

# Bivalent states



- Recall agents chose value 9 in last example
- But a network outage could look like agent 2 failing
- If fault-tolerant, Agents 1 & 3 might decide to output 7

→ Once network back, Agent 2 must also output 7

## Definition (Bivalent)

An execution of a consensus protocol is in a **bivalent** state when the network can affect which value agents choose.



# Univalent and stuck states

## Definition (Univalent, Valent)

An execution of a consensus protocol is in a **univalent** state when only one output value is possible. If that value is  $i$ , call the state  **$i$ -valent**.

## Definition (Stuck)

An execution of a [broken] consensus protocol is in a **stuck** state when one or more non-faulty nodes can never output a value.

- **Recall output is write once and all outputs must agree**
  - Hence, no output is possible in bivalent state
- **If an execution starts in a bivalent state and terminates, it must at some point reach a univalent state**

# FLP intuition

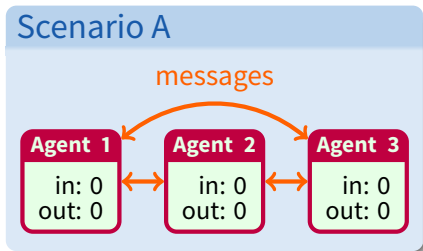
- Consider a terminating execution of a bivalent system
- Let  $m$  be last message received in a bivalent state
  - Call  $m$  the execution's **deciding message**
  - Any terminating execution requires a deciding message
- Suppose the network had delayed  $m$ 
  - Other messages could cause transitions to other bivalent states
  - Then, receiving  $m$  might no longer lead to a univalent state
  - In this case, we say  $m$  has been **neutralized**

## Overview of FLP proof.

1. There are bivalent starting configurations
2. The network can neutralize any deciding message
3. Hence, the system can remain bivalent in perpetuity



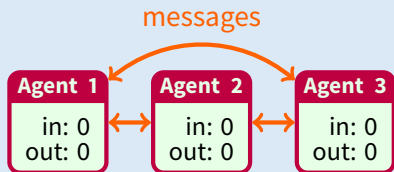
# There exists a bivalent state



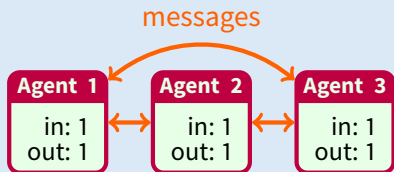
- Assume you could have liveness with an agent failure
- If all inputs 0, correct agents must eventually output 0
  - Similarly, if all inputs 1, correct agents must eventually output 1
- Now say we start flipping one input bit at a time
- Find 0- and 1-valent states differing at only one input
  - Suppose node with this differing input fails
  - By assumption, the system nonetheless reaches consensus
  - Hence output depends on network; at least one state was bivalent

# There exists a bivalent state

Scenario A

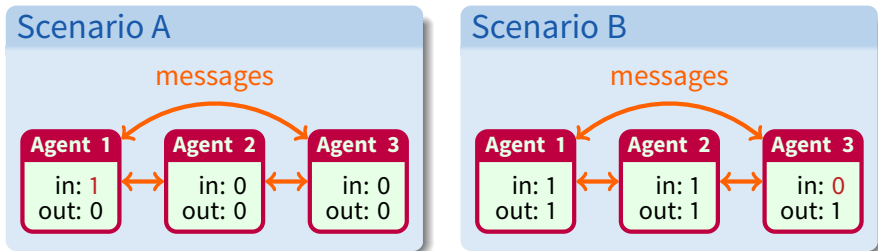


Scenario B



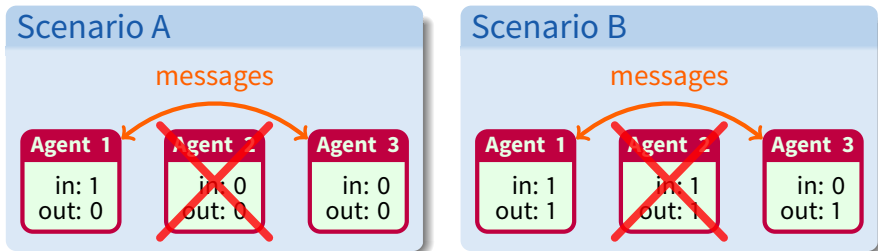
- Assume you could have liveness with an agent failure
- If all inputs 0, correct agents must eventually output 0  
→ Similarly, if all inputs 1, correct agents must eventually output 1
- Now say we start flipping one input bit at a time
- Find 0- and 1-valent states differing at only one input
  - Suppose node with this differing input fails
  - By assumption, the system nonetheless reaches consensus
  - Hence output depends on network; at least one state was bivalent

# There exists a bivalent state




- Assume you could have liveness with an agent failure
  - If all inputs 0, correct agents must eventually output 0
    - Similarly, if all inputs 1, correct agents must eventually output 1
  - Now say we start flipping one input bit at a time
- Find 0- and 1-valent states differing at only one input
- Suppose node with this differing input fails
  - By assumption, the system nonetheless reaches consensus
  - Hence output depends on network; at least one state was bivalent

# There exists a bivalent state




- Assume you could have liveness with an agent failure
- If all inputs 0, correct agents must eventually output 0
  - Similarly, if all inputs 1, correct agents must eventually output 1
- Now say we start flipping one input bit at a time
- Find 0- and 1-valent states differing at only one input
  - Suppose node with this differing input fails
    - By assumption, the system nonetheless reaches consensus
    - Hence output depends on network; at least one state was bivalent

# Any message can be neutralized



- Let  $m$  be a deciding message for value 0 from state  $b$
- Consider a message schedule from  $b$  to a 1-valent state
    - If  $m$  is on the path, it leads to a bi-valent state
    - If  $m$  is not on the path, append it to the (1-valent) path
  - Apply  $m$  to each node on the path
    - Either  $m$  will lead to a bi-valent state, or it will produce differing univalent states on adjacent nodes  $c_0$  and  $c_1$


# Any message can be neutralized



- Let  $m$  be a deciding message for value 0 from state  $b$
- Consider a message schedule from  $b$  to a 1-valent state
  - If  $m$  is on the path, it leads to a bi-valent state
  - If  $m$  is not on the path, append it to the (1-valent) path
- Apply  $m$  to each node on the path
  - Either  $m$  will lead to a bi-valent state, or it will produce differing univalent states on adjacent nodes  $c_0$  and  $c_1$




# Any message can be neutralized




- Let  $m$  be a deciding message for value 0 from state  $b$
- Consider a message schedule from  $b$  to a 1-valent state
  - If  $m$  is on the path, it leads to a bi-valent state
    - If  $m$  is not on the path, append it to the (1-valent) path
- Apply  $m$  to each node on the path
  - Either  $m$  will lead to a bi-valent state, or it will produce differing univalent states on adjacent nodes  $c_0$  and  $c_1$

# Any message can be neutralized



- Let  $m$  be a deciding message for value 0 from state  $b$
- Consider a message schedule from  $b$  to a 1-valent state
  - If  $m$  is on the path, it leads to a bi-valent state or to a 1-valent one
- If  $m$  is not on the path, append it to the (1-valent) path
- Apply  $m$  to each node on the path
  - Either  $m$  will lead to a bi-valent state, or it will produce differing univalent states on adjacent nodes  $c_0$  and  $c_1$

# Any message can be neutralized




- Let  $m$  be a deciding message for value 0 from state  $b$
- Consider a message schedule from  $b$  to a 1-valent state
  - If  $m$  is on the path, it leads to a bi-valent state or to a 1-valent one
  - If  $m$  is not on the path, append it to the (1-valent) path

→ **Apply  $m$  to each node on the path**


- Either  $m$  will lead to a bi-valent state, or it will produce differing univalent states on adjacent nodes  $c_0$  and  $c_1$

# Any message can be neutralized




- Let  $m'$  be the message that transitions between  $c_0$  and  $c_1$
- If  $m, m'$  received by different agents, order won't matter
    - But if delivering *both* messages yields a 1-valent state, delivering just  $m$  can't yield a 0-valent state
  - Hence, either  $m$  is neutralized at  $c_1$ , or same agent  $A$  received  $m$  and  $m'$ , making order significant
  - Yet if  $A$  slow after  $c_0$ , system must terminate without it

# Any message can be neutralized



- Let  $m'$  be the message that transitions between  $c_0$  and  $c_1$
- If  $m, m'$  received by different agents, order won't matter
  - But if delivering *both* messages yields a 1-valent state, delivering just  $m$  can't yield a 0-valent state
- Hence, either  $m$  is neutralized at  $c_1$ , or same agent  $A$  received  $m$  and  $m'$ , making order significant
- Yet if  $A$  slow after  $c_0$ , system must terminate without it


# Any message can be neutralized



→ Consider a run that terminates without  $A$


- Let  $x_1, \dots, x_n$  be the messages received (by nodes other than  $A$ )
- Let  $e$  be a univalent state reached during the run
- **Deliver  $x_1, \dots, x_n$  to terminating states after  $m$** 
  - Since  $ms$  and  $xs$  received by different nodes, can re-order
  - Means  $e$  not univalent (leads to both 0- and 1-valent states)!
- **Contradiction means  $m$  must be neutralized somewhere**

# Any message can be neutralized



- Consider a run that terminates without  $A$ 
  - Let  $x_1, \dots, x_n$  be the messages received (by nodes other than  $A$ )
  - Let  $e$  be a univalent state reached during the run
- Deliver  $x_1, \dots, x_n$  to terminating states after  $m$ 
  - Since  $m$ s and  $x$ s received by different nodes, can re-order
  - Means  $e$  not univalent (leads to both 0- and 1-valent states)!
- Contradiction means  $m$  must be neutralized somewhere

# Any message can be neutralized



- **Consider a run that terminates without  $A$** 
  - Let  $x_1, \dots, x_n$  be the messages received (by nodes other than  $A$ )
  - Let  $e$  be a univalent state reached during the run
- **Deliver  $x_1, \dots, x_n$  to terminating states after  $m$** 
  - Since  $m$ s and  $x$ s received by different nodes, can re-order

→ Means  $e$  not univalent (leads to both 0- and 1-valent states)!
- **Contradiction means  $m$  must be neutralized somewhere**



# Coping with FLP

- **This class will cover**
  - Many systems that require consensus
  - Many techniques for consensus
- **Safety is generally pretty important**
- **But can reasonably weaken liveness requirement**
  - Termination not guaranteed doesn't mean it won't happen
  - If your algorithm prevents completely stuck states  
...can often make it terminate “in practice”
- **Can weaken asynchronous system assumption**
- **Can make agents non-deterministic**
  - Have all nodes flip a coin to pick value—might all pick same value
  - Make it intractable for network to “guess” pathological delivery  
100% accurately in perpetuity